

Introduction to Python Data Analysis

Stephen Weston Robert Bjornson

Yale Center for Research Computing
Yale University

January 2017



Python for data analysis

Python is more of a general purpose programming language than R or Matlab. It has gradually become more popular for data analysis and scientific computing, but additional modules are needed. Some of the more popular modules are:

NumPy N-dimensional array

SciPy Scientific computing (linear algebra, numerical integration, optimization, etc)

Matplotlib 2D Plotting (similar to Matlab)

IPython Enhanced Interactive Console

Sympy Symbolic mathematics

Pandas Data analysis (provides a data frame structure similar to R)

NumPy, SciPy and Matplotlib are used in this presentation.

Creating N-dimensional arrays using NumPy

There are many ways to create N-dimensional arrays

```
import numpy as np

# Create 2X3 double precision array initialized to all zeroes
a = np.zeros((2,3), dtype=np.float64)

# Create array initialized by list of lists
a = np.array([[0,1,2],[3,4,5]], dtype=np.float64)

# Create array by reading CSV file
a = np.genfromtxt('data.csv', dtype=np.float64, delimiter=',')

# Create array using "arange" function
a = np.arange(6, dtype=np.float64).reshape(2,3)
```

Get values from N-dimensional array

NumPy provides many ways to extract data from arrays

Print single element of 2D array

```
print a[0,0]      # a scalar, not an array
```

Print first row of 2D array

```
print a[0,:]      # 1D array
```

Print last column of array

```
print a[:,-1]     # 1D array
```

Print sub-matrix of 2D array

```
print a[0:2,1:3]  # 2D array
```

Modifying N-dimensional arrays

NumPy uses the same basic syntax for modifying arrays

```
# Assign single value to single element of 2D array  
a[0,0] = 25.0
```

```
# Assign 1D array to first row of 2D array  
a[0,:] = np.array([10,11,12], dtype=np.float64)
```

```
# Assign 1D array to last column of 2D array  
a[:,-1] = np.array([20,21], dtype=np.float64)
```

```
# Assign 2D array to sub-matrix of 2D array  
a[0:2,1:3] = np.array([[10,11],[20,21]], dtype=np.float64)
```

Modifying arrays using broadcasting

```
# Assign scalar to first row of 2D array
```

```
a[0,:] = 10.0
```

```
# Assign 1D array to all rows of 2D array
```

```
a[:,:] = np.array([30,31,32], dtype=np.float64)
```

```
# Assign 1D array to all columns of 2D array
```

```
a[:,:] = np.array([40,41], dtype=np.float64).reshape(2,1)
```

```
# Assign scalar to sub-matrix of 2D array
```

```
a[0:2,1:3] = 100.0
```

Arithmetic on arrays

Operate on arrays using binary operators and NumPy functions

```
# Create 1D array
```

```
a = np.arange(4, dtype=np.float64)
```

```
# Add 1D arrays elementwise
```

```
a + a
```

```
# Multiply 1D arrays elementwise
```

```
a * a
```

```
# Sum elements of 1D array
```

```
a.sum()
```

```
# Compute dot product
```

```
np.dot(a, a)      # same as: (a * a).sum()
```

```
# Compute cross product
```

```
np.dot(a.reshape(4,1), a.reshape(1,4))
```

NumPy views

Views are arrays that share memory with another array.

- views can make your program more memory and CPU efficient
- views are explicitly generated via the view method
- reshape and transpose implicitly return views of the original array
- arrays generated by slicing are views of the original
- use the copy method to avoid sharing memory
- set the writeable flag to make a view read-only (`a.flags.writeable`)

NumPy views continued

```
>>> a = np.arange(10)
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([ 0,  1, 100,  3,  4,  5,  6,  7,  8,  9])
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([  0,   1, 100,   3,   4,   5,   6,   7,   8,   9])
>>> a.__array_interface__['data'][0]
4330625024
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([  0,   1, 100,   3,   4,   5,   6,   7,   8,   9])
>>> a.__array_interface__['data'][0]
4330625024
>>> b.__array_interface__['data'][0]
4330625040
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([  0,   1, 100,   3,   4,   5,   6,   7,   8,   9])
>>> a.__array_interface__['data'][0]
4330625024
>>> b.__array_interface__['data'][0]
4330625040
>>> c = b.copy()
```


NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([  0,   1, 100,   3,   4,   5,   6,   7,   8,   9])
>>> a.__array_interface__['data'][0]
4330625024
>>> b.__array_interface__['data'][0]
4330625040
>>> c = b.copy()
>>> c.flags.owndata, c.__array_interface__['data'][0]
(True, 4301585776)
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([  0,   1, 100,   3,   4,   5,   6,   7,   8,   9])
>>> a.__array_interface__['data'][0]
4330625024
>>> b.__array_interface__['data'][0]
4330625040
>>> c = b.copy()
>>> c.flags.owndata, c.__array_interface__['data'][0]
(True, 4301585776)
>>> d = a
```

NumPy views continued

```
>>> a = np.arange(10)
>>> b = a[2::2]
>>> a.flags.owndata, b.flags.owndata
(True, False)
>>> b[0] = 100
>>> a
array([  0,   1, 100,   3,   4,   5,   6,   7,   8,   9])
>>> a.__array_interface__['data'][0]
4330625024
>>> b.__array_interface__['data'][0]
4330625040
>>> c = b.copy()
>>> c.flags.owndata, c.__array_interface__['data'][0]
(True, 4301585776)
>>> d = a
>>> d is a
True
```

Mini NumPy Cookbook

<i>matrix size</i>	<code>a.shape</code>
<i>transpose</i>	<code>a.T</code>
<i>extract diagonal elements</i>	<code>np.diag(a)</code>
<i>matrix multiply</i>	<code>a.dot(b)</code>
<i>element-wise multiply</i>	<code>a * b</code>
<i>column sums</i>	<code>np.sum(a, axis=0)</code>
<i>row sums</i>	<code>np.sum(a, axis=1)</code>
<i>element sum</i>	<code>np.sum(a, axis=None)</code>
<i>column means</i>	<code>np.mean(a, axis=0)</code>
<i>column max</i>	<code>np.max(a, axis=0)</code>
<i>column min</i>	<code>np.min(a, axis=0)</code>
<i>element-wise maximum</i>	<code>np.maximum(a, b)</code>
<i>element-wise minimum</i>	<code>np.minimum(a, b)</code>
<i>identity matrix</i>	<code>np.eye(n)</code>
<i>random matrix</i>	<code>np.random.rand(m, n)</code>
<i>set random seed</i>	<code>np.random.seed(i)</code>
<i>solve $ax=b$</i>	<code>np.linalg.solve(a, b)</code>
<i>fourier transform</i>	<code>np.fft.fft(a)</code>

SciPy Linear Algebra functions

```
import numpy as np
from scipy import linalg
a = np.array([[1, 2], [3, 4]], dtype=np.float64)

# Compute the inverse matrix
linalg.inv(a)

# Compute singular value decomposition
linalg.svd(a)

# Compute eigenvalues
linalg.eigvals(a)
```

2D plotting using Matplotlib

```
import numpy as np
import matplotlib.pyplot as plt
x = np.linspace(0.0, 2.0, 20)

plt.plot(x, np.sqrt(x), 'ro') # red circles
plt.show()

plt.plot(x, np.sqrt(x), 'b-') # blue lines
plt.show()

# Three plots in one figure
plt.plot(x, x, 'g--', x, np.sqrt(x), 'ro', x, np.sqrt(x), 'b-')
plt.show()
```